

CROATRAD

February 21, 2025

1 Analysis of CROATRAD.TXT from Statistics and Data Analysis in Geology by John C. Davis

In the book, the dataset has been used in Page No. 30 (Fig. 2.11) to show different histograms

1.1 What does the dataset contain?

The dataset contains measurement of absorbed background radiation levels per hour (unit: nGy/hr) in the Istrian Peninsula, Croatia. The radiation levels are measured for ^{40}K , ^{224}Ra , ^{228}Ra , ^{236}U and ^{137}Cs .

1. What is a Gray (Gy)?

- Gray (Gy) is the SI unit of absorbed radiation dose.
- It represents the amount of energy absorbed per kilogram of material.
- 1 Gy = 1 joule of radiation energy absorbed per kilogram of matter.

2. What is a Nanogray (nGy)?

- 1 Nanogray (nGy) = 10⁻⁹ Gray (Gy) = one billionth of a Gray.
- Since natural background radiation levels are very low, we use nGy instead of Gy.

3. What Does “per Hour” (nGy/hr) Mean?

- nGy/hr (Nanogray per Hour) measures the radiation dose rate.
- It tells how much radiation energy is absorbed every hour.

The last two columns show the Natural and the Total Radiation levels in the Istrian Peninsula.

Comparison with Safety Limits for Human Exposure

Source	Radiation Level	Comparison
Typical background radiation	50-200 nGy/hr	Matches most of the data
ICRP Public Exposure Limit (International Commission on Radiological Protection)	~1,000 nGy/hr (1 μ Gy/hr)	Much higher than the observed values
UNSCEAR (United Nations Scientific Committee on the Effects of Atomic Radiation) natural range	400-1,500 nGy/hr	Istria's values are well below concern levels
Chernobyl-affected areas today	5,000-20,000 nGy/hr	Much higher than Istrian levels

1.2 Load Necessary Python Modules

```
[1]: import pandas as pd      # for file handling and manipulation of dataframes  
[2]: import numpy as np       # for handling arrays and for advanced histograms in this chapter  
[3]: import matplotlib.pyplot as plt    # for data visualization
```

To view the plots within the IPython cell in Jupyter Notebook:

```
[4]: %matplotlib inline  
[5]: import seaborn as sns     # for statistical visualizations not covered by matplotlib
```

1.3 Read the data and convert it into Data Frame in Python using pandas

```
[6]: # Step 1: Load the Data in Python  
file_path = "CROATRAD.TXT"  
  
[7]: # Create dataframe and convert the .txt file to csv using the first row of the .txt file as the headers  
df = pd.read_csv(file_path, delim_whitespace=True, names=["40K", "224Ra", "228Ra", "236U", "137Cs", "Natural", "Total"])  
df.head()
```

C:\Users\DELL\AppData\Local\Temp\ipykernel_28308\2950647005.py:2: FutureWarning:
The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed
in a future version. Use ``sep='\\s+'`` instead
df = pd.read_csv(file_path, delim_whitespace=True, names=["40K", "224Ra", "228Ra", "236U", "137Cs", "Natural", "Total"])

```
[7]:   40K  224Ra  228Ra  236U  137Cs  Natural  Total  
0   40K  224Ra  228Ra  236U  137Cs  Natural  Total  
1  13.7   41.9   24.4   0.1    2.8    80.2    83  
2   14    52.6   48.4   0.4    3.3   115.4  118.7  
3  12.3   40.6   27.7   0.2    1.6    80.8   82.3  
4  13.2   45.9   29.6   0.4    2.5    89.2   91.6
```

Since the above code is being deprecated, the following code can be used:

```
df = pd.read_csv(file_path, sep=r'\s+', names=["40K", "224Ra", "228Ra", "236U", "137Cs", "Natural", "Total"])  
  
[8]: df = pd.read_csv(file_path, sep=r'\s+', names=["40K", "224Ra", "228Ra", "236U", "137Cs", "Natural", "Total"])  
df.head()
```

```
[8]:    40K  224Ra  228Ra  236U  137Cs  Natural  Total
      0   40K  224Ra  228Ra  236U  137Cs  Natural  Total
      1  13.7   41.9   24.4   0.1    2.8    80.2    83
      2   14.0   52.6   48.4   0.4    3.3   115.4  118.7
      3  12.3   40.6   27.7   0.2    1.6    80.8   82.3
      4  13.2   45.9   29.6   0.4    2.5    89.2   91.6
```

However, in the output, the first row (indexed 0) is the same as the header. Therefore, the best code would be:

```
df = pd.read_csv(file_path, sep=r'\s+', header = 0)
```

```
[9]: df = pd.read_csv(file_path, sep=r'\s+', header = 0)
df
```

```
[9]:    40K  224Ra  228Ra  236U  137Cs  Natural  Total
      0  13.7   41.9   24.4   0.1    2.8    80.2   83.0
      1  14.0   52.6   48.4   0.4    3.3   115.4  118.7
      2  12.3   40.6   27.7   0.2    1.6    80.8   82.3
      3  13.2   45.9   29.6   0.4    2.5    89.2   91.6
      4  11.3   40.6   21.6   0.1    1.6    73.6   75.2
      ..
      ..
      ..
      120 15.4   22.6   12.7   0.1    4.7    50.8   55.6
      121 16.7   14.6   8.5    0.1    3.4    39.9   43.2
      122 15.6   14.6   8.0    0.2    1.2    38.4   39.6
      123 18.7   47.3   35.7   0.2    1.8   101.9  103.7
      124 13.3   40.6   31.0   0.2    0.8    85.2   85.9
```

[125 rows x 7 columns]

1.4 Measurement of Central Tendencies using Python

1.4.1 Means

```
[10]: print(df["40K"].mean())
```

15.0224

```
[11]: print(df["224Ra"].mean())
```

36.14160000000001

```
[12]: print(df["228Ra"].mean())
```

36.8016

```
[13]: print(df["236U"].mean())
```

0.2848000000000005

```
[14]: print(f"{df['236U'].mean():.4f}")
```

```
0.2848
```

```
[15]: print(df["137Cs"].mean())
```

```
6.177599999999999
```

```
[16]: print(f"{df['137Cs'].mean():.4f}")
```

```
6.1776
```

```
[17]: print(df["Natural"].mean())
```

```
88.2519999999998
```

```
[18]: print(f"{df['Natural'].mean():.4f}")
```

```
88.2520
```

```
[19]: print(df["Total"].mean())
```

```
94.4184
```

1.4.2 Medians

```
[20]: print(df["40K"].median())
```

```
15.0
```

```
[21]: print(df["224Ra"].median())
```

```
39.3
```

```
[22]: print(df["228Ra"].median())
```

```
31.5
```

```
[23]: print(df["236U"].median())
```

```
0.2
```

```
[24]: print(df["137Cs"].median())
```

```
4.4
```

```
[25]: print(df["Natural"].median())
```

```
90.8
```

```
[26]: print(df["Total"].median())
```

```
96.0
```

1.4.3 Simpler way of Getting Descriptive Statistics

It's rather tedious to understand the descriptive statistics of the data in the way given above. Pandas provides a better solution. We can use:

DataFrame.mean() to find the mean of each column

DataFrame.median() to find the median of each column

...

We can also type:

DataFrame.describe() to get the descriptive statistics in one go.

[27]: df.mean()

```
[27]: 40K      15.0224
224Ra    36.1416
228Ra    36.8016
236U     0.2848
137Cs    6.1776
Natural   88.2520
Total     94.4184
dtype: float64
```

[28]: df.median()

```
[28]: 40K      15.0
224Ra    39.3
228Ra    31.5
236U     0.2
137Cs    4.4
Natural   90.8
Total     96.0
dtype: float64
```

[29]: df.describe()

```
[29]:      40K      224Ra      228Ra      236U      137Cs    Natural \
count  125.000000  125.000000  125.000000  125.000000  125.000000  125.000000
mean   15.022400  36.141600  36.801600  0.284800   6.177600  88.252000
std    2.750304  11.922405  24.065301  0.180079   5.999541  31.467419
min    5.200000  8.000000  7.100000  0.000000   0.600000  31.600000
25%   13.200000  26.600000  21.600000  0.200000   2.400000  66.800000
50%   15.000000  39.300000  31.500000  0.200000   4.400000  90.800000
75%   16.700000  44.600000  47.000000  0.300000   7.500000  105.400000
max   22.400000  59.200000 146.200000  1.000000  37.100000 211.800000

      Total
count  125.000000
```

```
mean      94.418400
std       32.470407
min      36.000000
25%     77.100000
50%     96.000000
75%    109.700000
max     218.200000
```

1.5 Histogram for Total Radiation Dose Absorbed Per Hour using Python

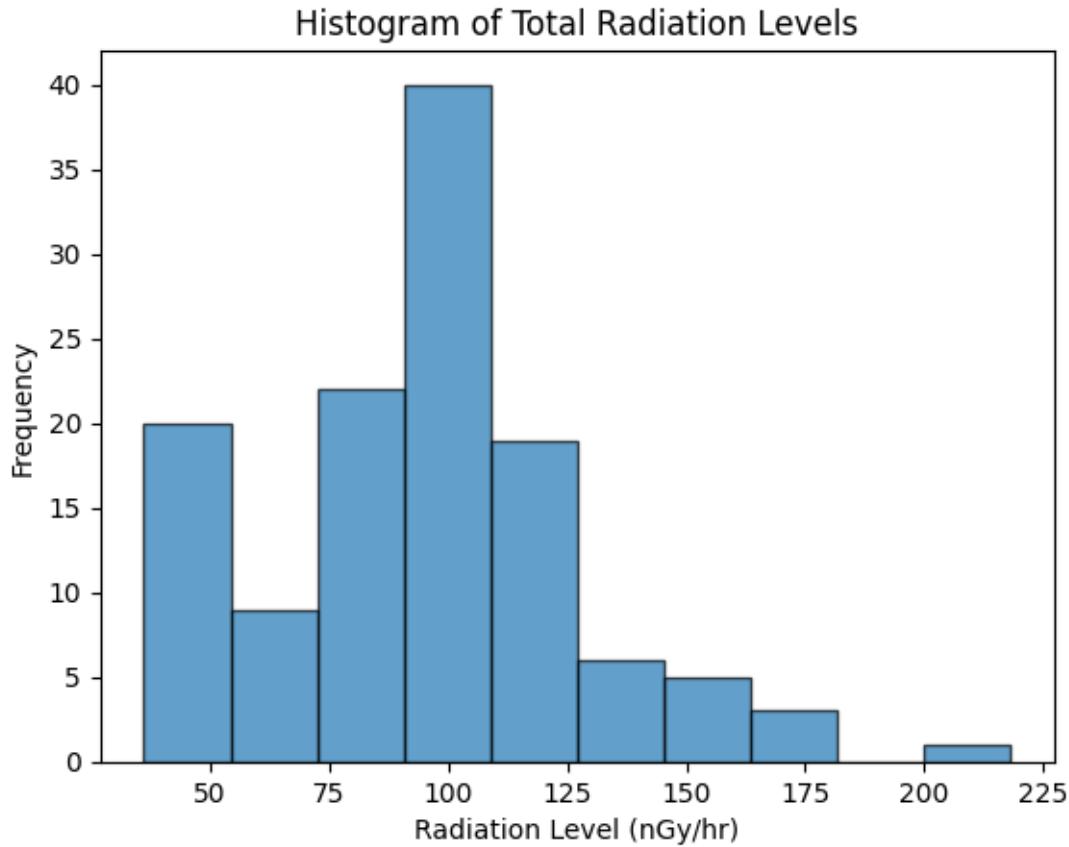
1.5.1 Basic Histogram

Using matplotlib to visualize the distribution of Total Radiation.

```
[30]: # histogram
plt.hist(df["Total"], bins = 10, edgecolor = "black", alpha = 0.7)

# Labels and Title
plt.xlabel("Radiation Level (nGy/hr)")
plt.ylabel("Frequency")
plt.title("Histogram of Total Radiation Levels")

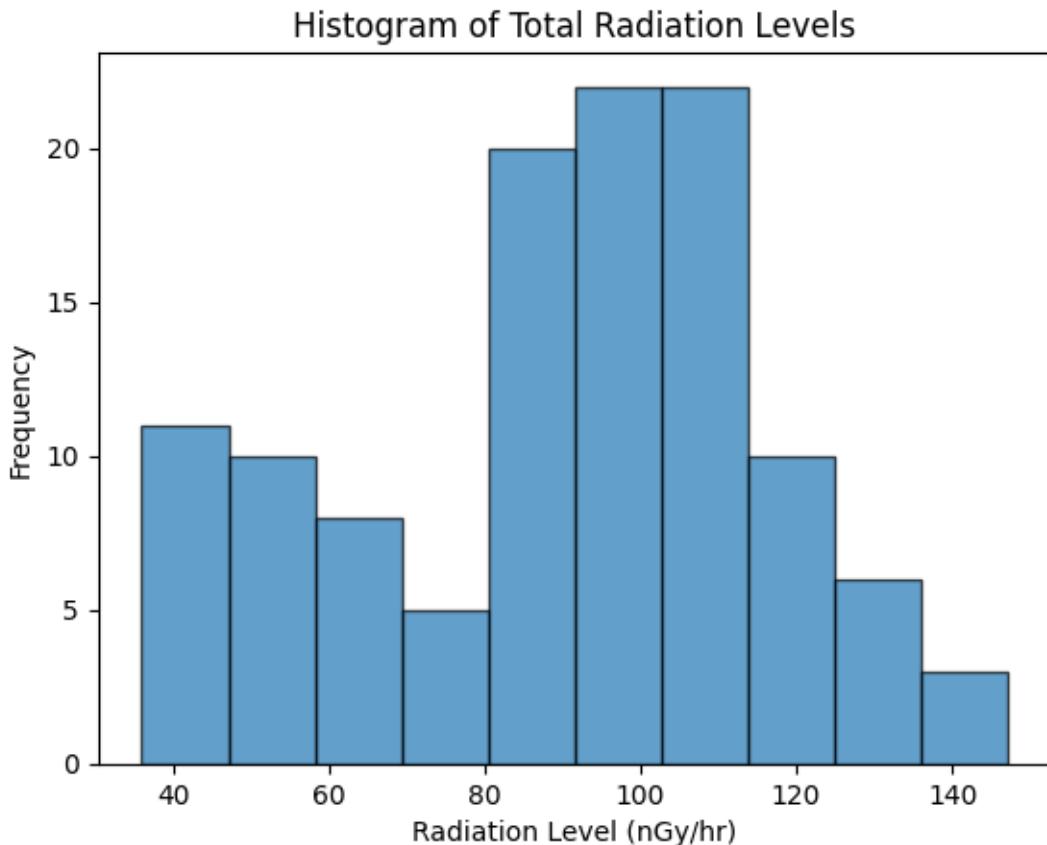
# Show the plot
plt.show()
```



Basic Histogram taking values < 150 nGy/hr

```
[31]: filtered_data = df[df["Total"] < 150]["Total"]
plt.hist(filtered_data, bins = 10, edgecolor = "black", alpha = 0.7)
# Labels and Title
plt.xlabel("Radiation Level (nGy/hr)")
plt.ylabel("Frequency")
plt.title("Histogram of Total Radiation Levels")

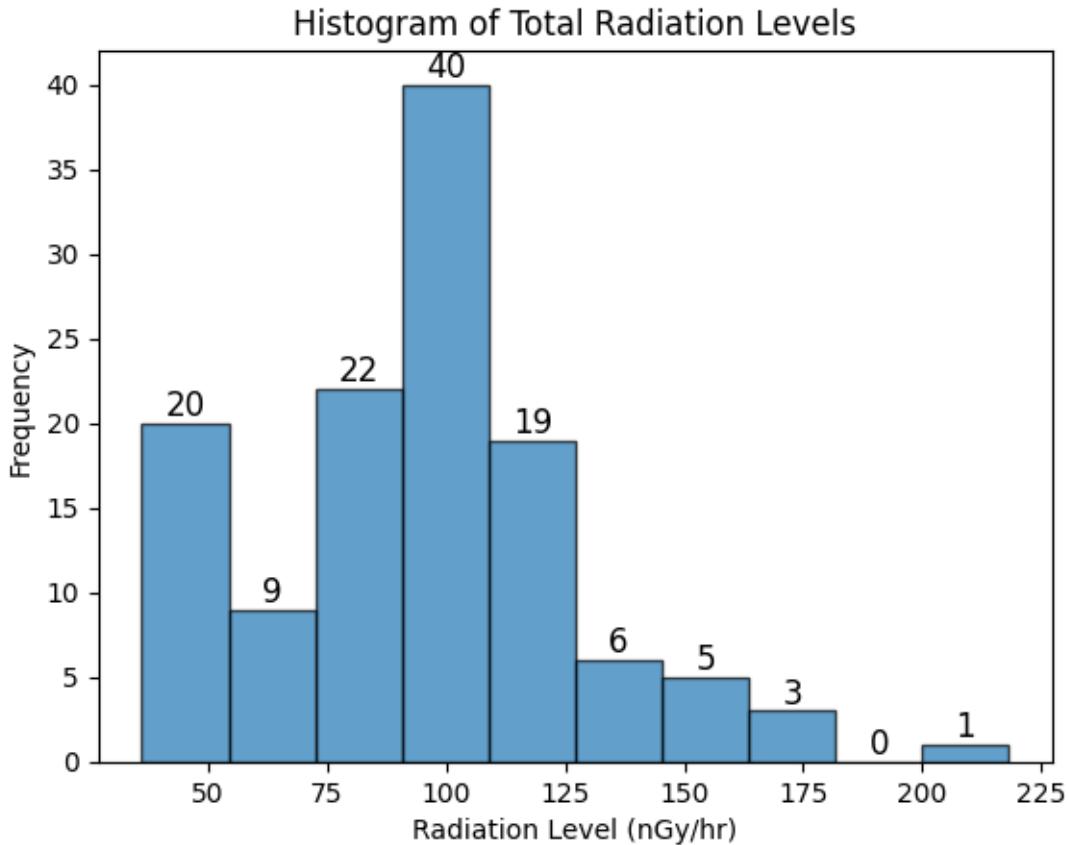
# Show the plot
plt.show()
```



```
[32]: counts, bin_edges, _ = plt.hist(df["Total"], bins=10, edgecolor="black", alpha=0.7)

# Add labels on top of bars
for count, bin_edge in zip(counts, bin_edges[:-1]):
    plt.text(bin_edge + (bin_edges[1] - bin_edges[0]) / 2, count,
             str(int(count)),
             ha='center', va='bottom', fontsize=12)
# Labels and Title
plt.xlabel("Radiation Level (nGy/hr)")
plt.ylabel("Frequency")
plt.title("Histogram of Total Radiation Levels")

# Show the plot
plt.show()
```



```
[33]: bin_width = bin_edges[1] - bin_edges[0]
print(bin_width)
```

18.22

```
[34]: # histogram
fig, ax = plt.subplots(figsize=(8, 6))
counts, bin_edges, patches = ax.hist(df["Total"], bins=10, edgecolor="black", alpha=0.7)

# Add labels on top of bars
for count, bin_edge in zip(counts, bin_edges[:-1]):
    ax.text(bin_edge + (bin_edges[1] - bin_edges[0]) / 2, count,
            str(int(count)),
            ha='center', va='bottom', fontsize=12)

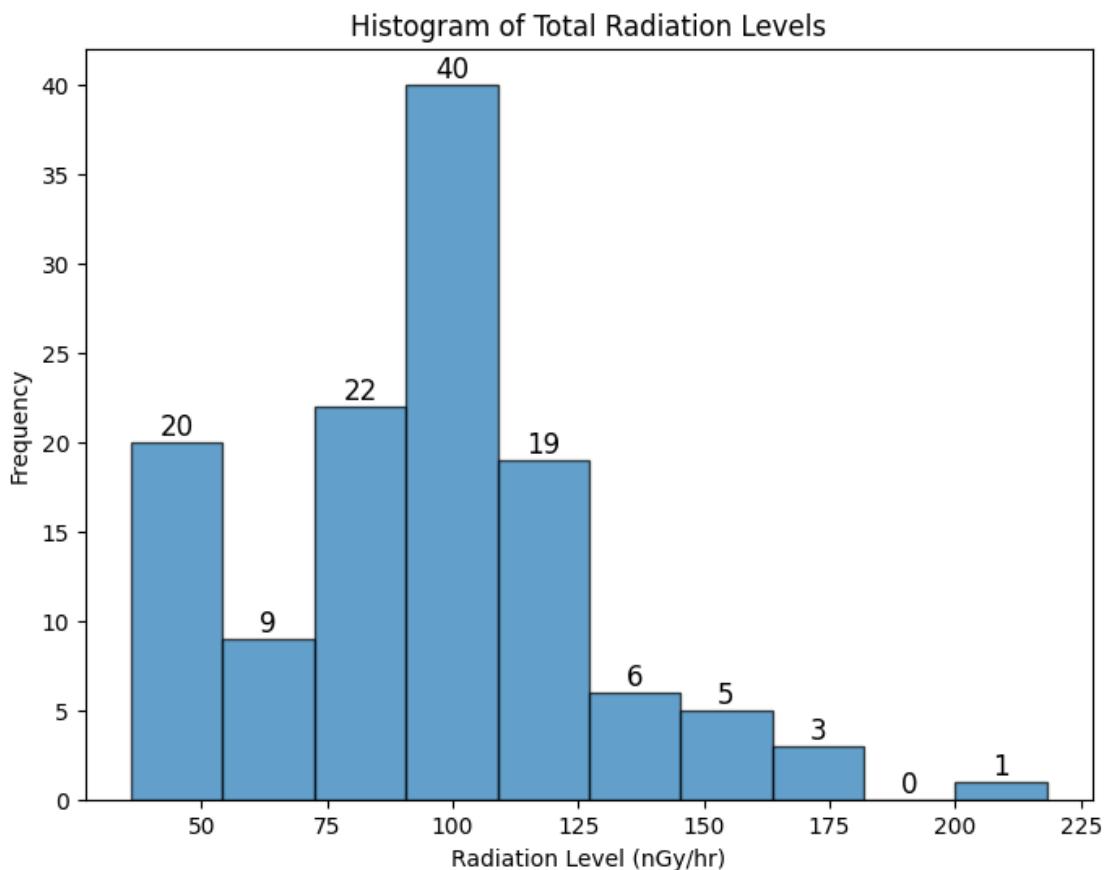
# Labels and Title
ax.set_xlabel("Radiation Level (nGy/hr)")
ax.set_ylabel("Frequency")
```

```

ax.set_title("Histogram of Total Radiation Levels")

# Show the plot
plt.show()

```



1.5.2 Histograms with different Bin Sizes

```

[35]: fig, axes = plt.subplots(2, 2, figsize = (10, 10))

# 5 bins
axes[0, 0].hist(df["Total"], bins = 5, edgecolor = "black", alpha = 0.7)
axes[0, 0].set_title("Histogram (5 bins)")

# 7 bins
axes[0, 1].hist(df["Total"], bins = 7, edgecolor = "black", alpha = 0.7)
axes[0, 1].set_title("Histogram (7 bins)")

# 10 bins
axes[1, 0].hist(df["Total"], bins = 10, edgecolor = "black", alpha = 0.7)

```

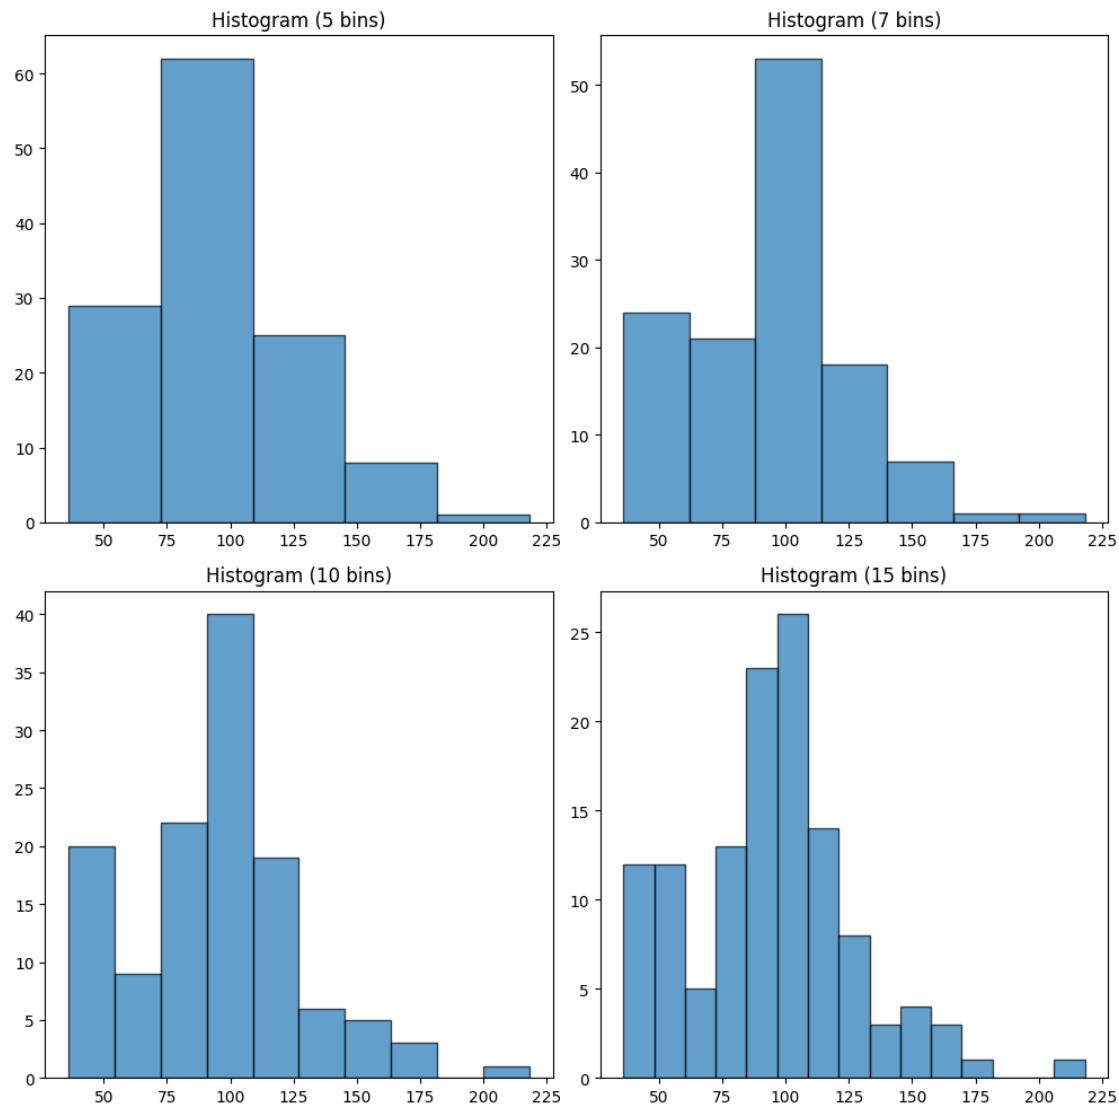
```

axes[1, 0].set_title("Histogram (5 bins)")

# 15 bins
axes[1, 1].hist(df["Total"], bins = 15, edgecolor = "black", alpha = 0.7)
axes[1, 1].set_title("Histogram (15 bins)")

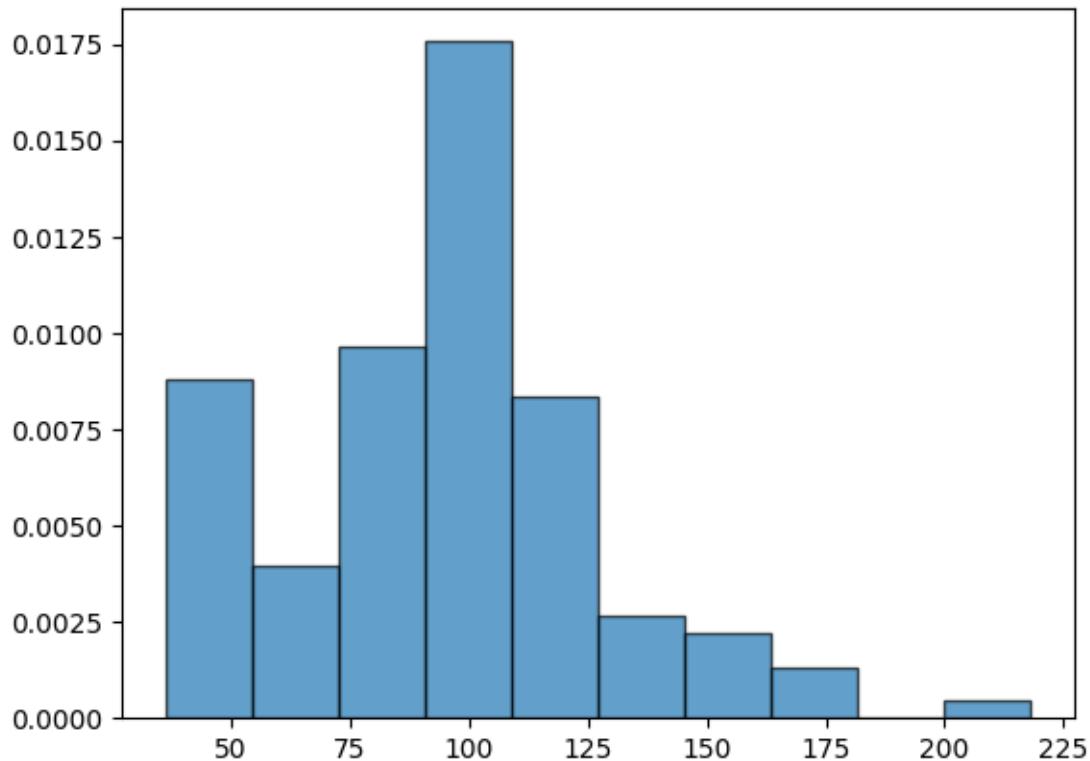
plt.tight_layout()
plt.show()

```



1.5.3 Normalized Histogram (Percentage)

```
[36]: plt.hist(df["Total"], bins = 10, density = True, edgecolor = "black", alpha = 0.7)
plt.show()
```

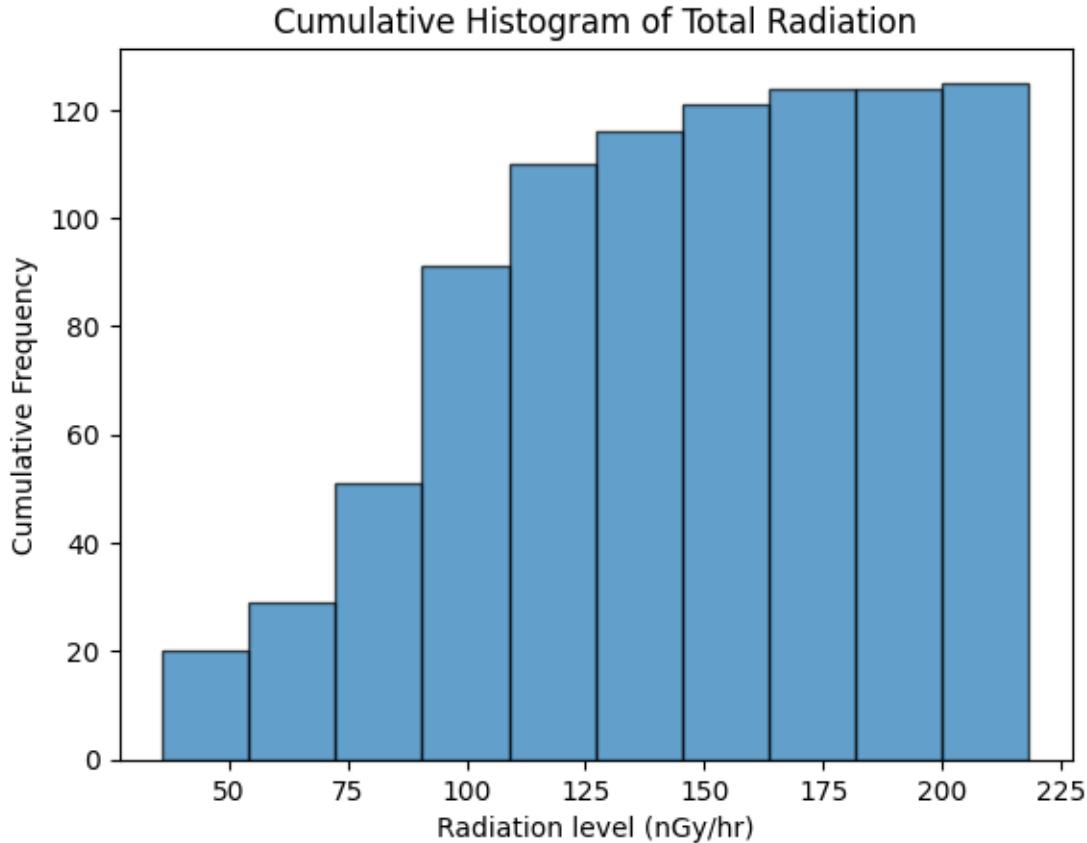


1.5.4 Cumulative Histogram

```
[37]: plt.hist(df["Total"], bins = 10, cumulative = True, edgecolor = "black", alpha= 0.7)

plt.xlabel("Radiation level (nGy/hr)")
plt.ylabel("Cumulative Frequency")
plt.title("Cumulative Histogram of Total Radiation")

plt.show()
```

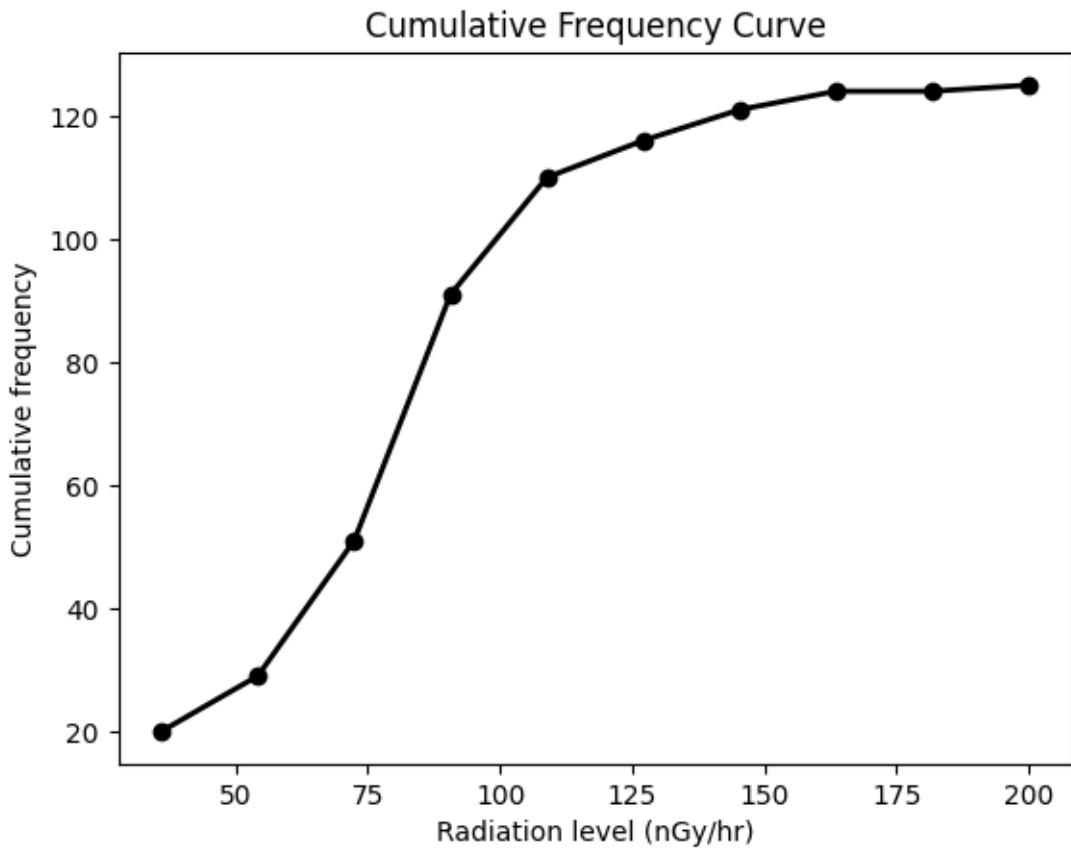


The cumulative histogram looks quite odd. It could be better represented with the cumulative frequency curve.

```
[38]: # Calculate cumulative frequency using numpy
counts, bin_edges = (np.histogram(df["Total"], bins = 10))
cumulative_counts = np.cumsum(counts)

# Plot cumulative frequency curve using matplotlib
plt.plot(bin_edges[:-1], cumulative_counts, linewidth = 2, linestyle = '-',
          color = "black", marker = 'o')
plt.xlabel("Radiation level (nGy/hr)")
plt.ylabel("Cumulative frequency")
plt.title("Cumulative Frequency Curve")

plt.show()
```



Cumulative frequency percentage curve

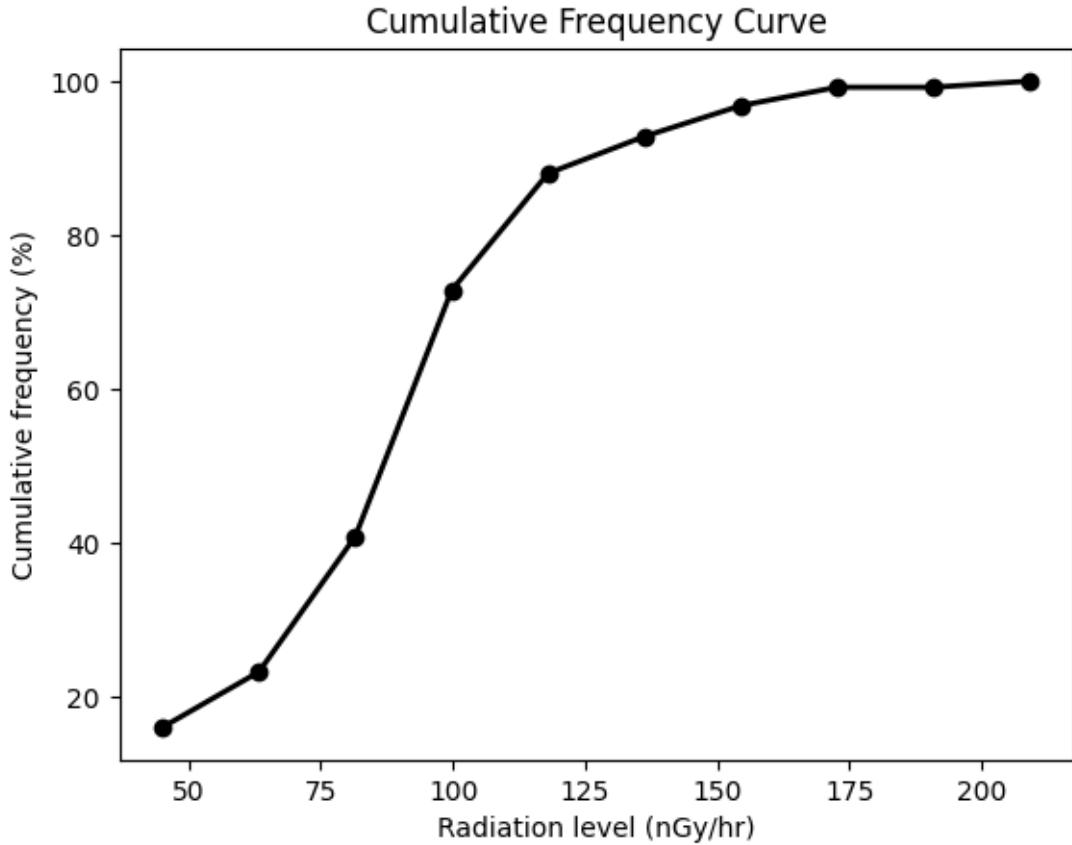
```
[39]: # Calculate cumulative frequency using numpy
counts, bin_edges = (np.histogram(df["Total"], bins = 10))

cumulative_counts = np.cumsum(counts)
cumulative_percentage = (cumulative_counts / cumulative_counts[-1]) * 100

bin_centers = (bin_edges[:-1] + bin_edges[1:]) / 2

# Plot cumulative frequency curve using matplotlib
plt.plot(bin_centers, cumulative_percentage, linewidth = 2, linestyle = '-',
          color = "black", marker = 'o')
plt.xlabel("Radiation level (nGy/hr)")
plt.ylabel("Cumulative frequency (%)")
plt.title("Cumulative Frequency Curve")

plt.show()
```



Basic Histogram and Cumulative Percentage Curve Combined in the Same Plot

```
[40]: # Calculate cumulative frequency using numpy
counts, bin_edges = (np.histogram(df["Total"], bins = 10))

# Calculate bin width for plotting histogram (bar plot)
bin_width = (bin_edges[1] - bin_edges[0])

# Calculate cumulative frequency and cumulative frequency percentage
cumulative_counts = np.cumsum(counts)
cumulative_percentage = (cumulative_counts / cumulative_counts[-1]) * 100

# Calculate bin centers for plotting cumulative frequency curve
bin_centers = (bin_edges[:-1] + bin_edges[1:]) / 2

# Create a dual-axis figure
fig, ax1 = plt.subplots(figsize = (8, 6))
ax2 = ax1.twinx()

# Plot histogram (frequency/count) on the Left axis
```

```

ax1.bar(bin_centers, counts, bin_width, alpha = 0.7, color = "cornflowerblue", □
    ↪edgecolor = "black", label = "Frequency")

# Set x-label and y-label
ax1.set_xlabel("Radiation level (nGy/hr)")
ax1.set_ylabel("Frequency")
ax1.tick_params(axis = "y", labelcolor = "cornflowerblue")

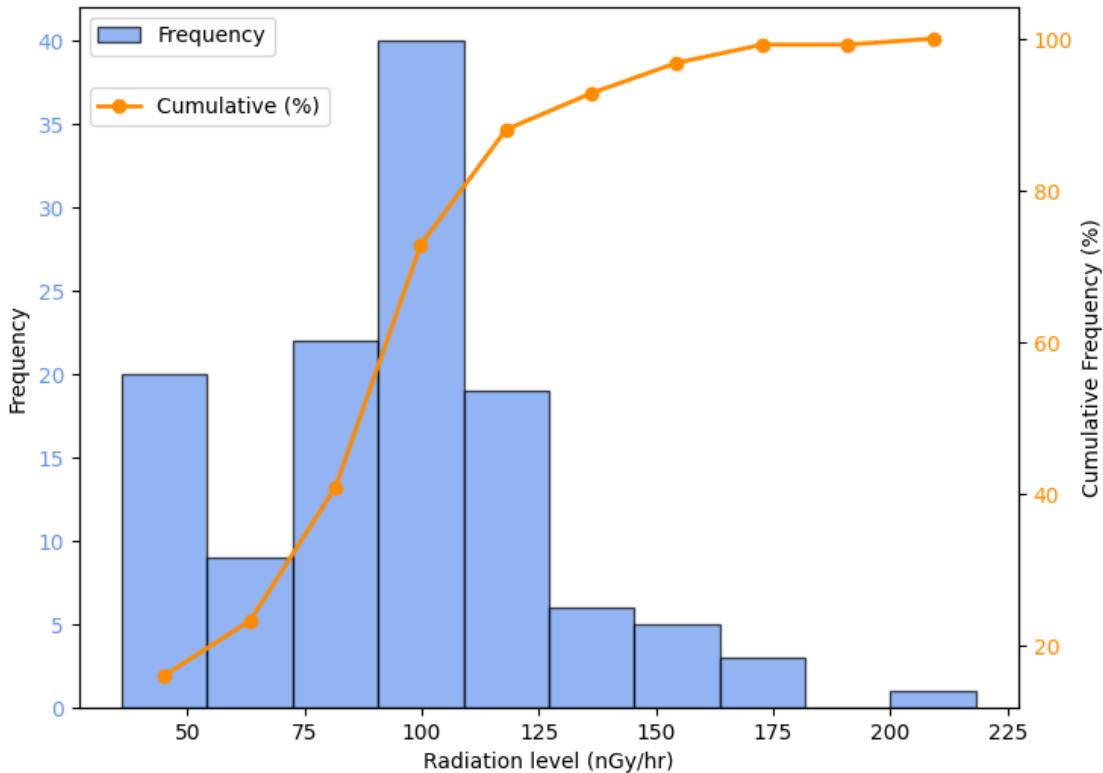
# Plot cumulative frequency curve on the Right axis
ax2.plot(bin_centers, cumulative_percentage, marker = "o", linestyle = "-",
    ↪linewidth = 2, color = "darkorange", label = "Cumulative (%)")
ax2.set_ylabel("Cumulative Frequency (%)")
ax2.tick_params(axis = "y", labelcolor = "darkorange")

# Title and Legend
fig.suptitle("Histogram with Cumulative Percentage", fontsize=14)
ax1.legend(loc='upper left') # Move frequency label to upper left
ax2.legend(loc='upper left', bbox_to_anchor=(0, 0.9)) # Move cumulative
    ↪percentage below frequency in left corner

# Show plot
plt.show()

```

Histogram with Cumulative Percentage

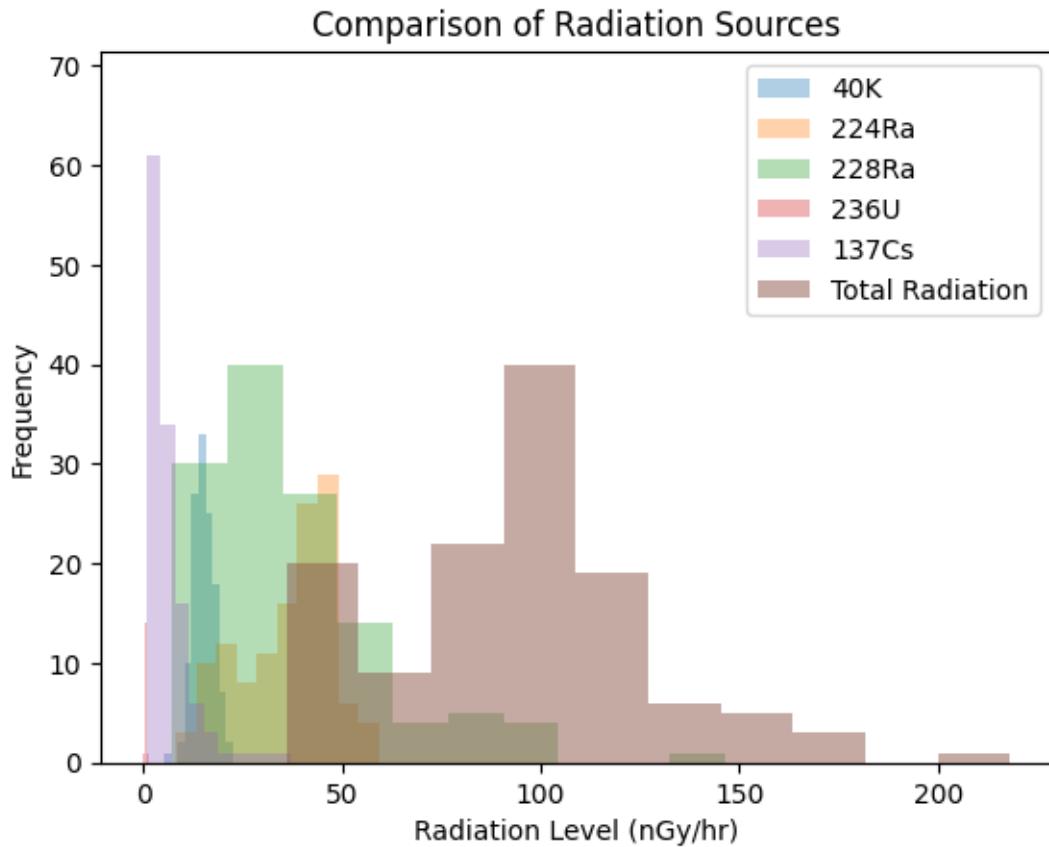


1.5.5 Multiple Histograms

```
[41]: plt.hist(df["40K"], bins=10, alpha=0.35, label="40K")
plt.hist(df["224Ra"], bins=10, alpha=0.35, label="224Ra")
plt.hist(df["228Ra"], bins=10, alpha=0.35, label="228Ra")
plt.hist(df["236U"], bins=10, alpha=0.35, label="236U")
plt.hist(df["137Cs"], bins=10, alpha=0.35, label="137Cs")
plt.hist(df["Total"], bins=10, alpha=0.5, label="Total Radiation")

plt.xlabel("Radiation Level (nGy/hr)")
plt.ylabel("Frequency")
plt.title("Comparison of Radiation Sources")
plt.legend()

plt.show()
```



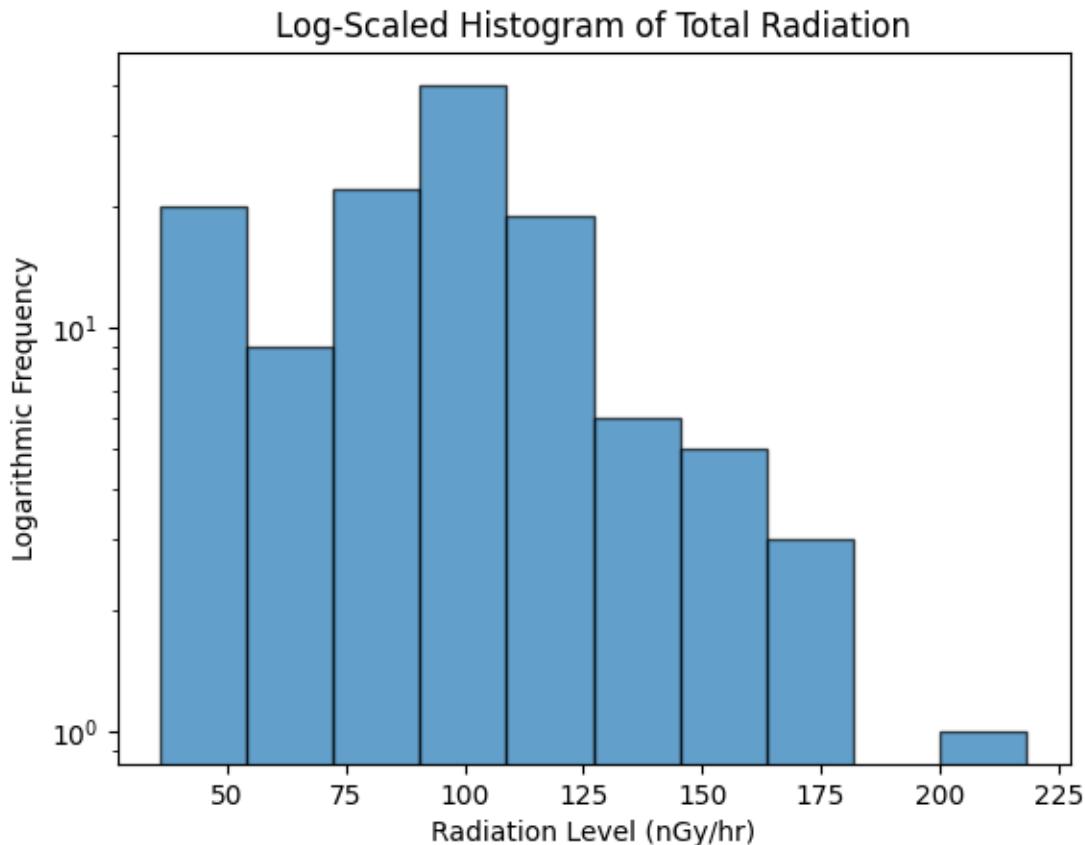
1.5.6 Logarithmic Scale Histogram

Just a demo, not meaningful for this data

```
[42]: plt.hist(df["Total"], bins=10, edgecolor="black", alpha=0.7)
plt.yscale("log") # Set Y-axis to logarithmic scale

plt.xlabel("Radiation Level (nGy/hr)")
plt.ylabel("Logarithmic Frequency")
plt.title("Log-Scaled Histogram of Total Radiation")

plt.show()
```

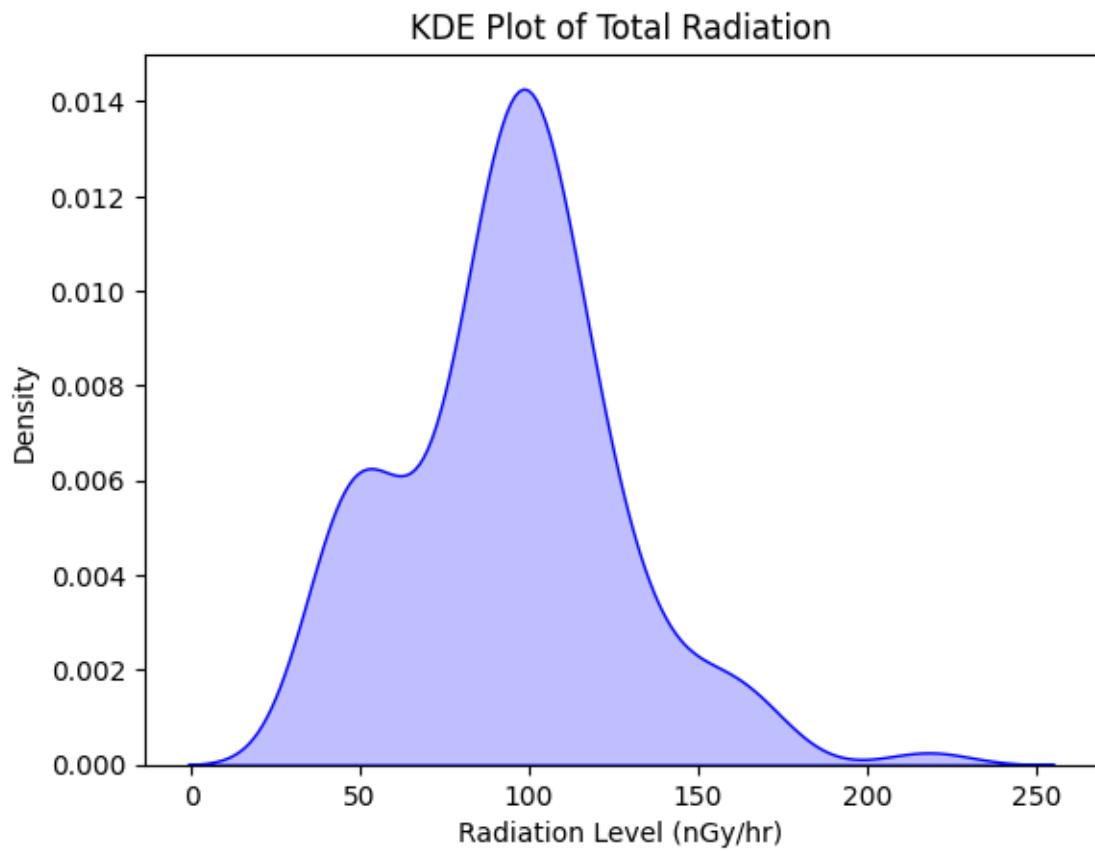


1.5.7 KDE (Kernel Density Estimation)/Smoothed Histogram Using Seaborn

```
[43]: sns.kdeplot(df["Total"], fill = True, color = "blue")

plt.xlabel("Radiation Level (nGy/hr)")
plt.ylabel("Density")
plt.title("KDE Plot of Total Radiation")
```

```
plt.show()
```



This is a continuous approximation of the distribution

1.5.8 Choosing the Appropriate Histogram?

1. Basic Histogram : General Overview of Distribution
2. Different bin sizes : To see how bin size affects interpretation
3. Normalized Histogram : To show percentage instead of counts
4. Cumulative Histogram : Understanding threshold exceedance
5. Multiple Overlaid Histograms : Comparing multiple data
6. Logarithmic Histogram : Data spanning large magnitude differences
7. KDE (Smoothed) : Continuous approximation of distribution

1.5.9 The Importance of Histograms in Data Analysis

A **histogram** is one of the most fundamental tools in data analysis, especially for visualizing distributions. Here's why histograms are so useful:

1. Understanding Data Distribution

Histograms show how data values are distributed across different ranges. This helps in identifying:

- **Central Tendency:** The most frequent values in a dataset.
- **Dispersion:** How spread out the values are.
- **Skewness:** Whether the data is symmetrical or skewed to one side.

Example: In radiation data analysis, histograms help reveal whether most locations have normal radiation levels or if some areas have **outliers** with unusually high exposure.

2. Identifying Patterns and Trends

Histograms make it easy to spot patterns in data:

- **Normal Distribution:** A bell-shaped histogram suggests that most values are clustered around the mean.
- **Bimodal or Multimodal Distributions:** Two or more peaks indicate multiple distinct groups within the data.
- **Uniform Distribution:** A flat histogram suggests equal frequency across the range.

Example: If radiation levels in Istria had two peaks (one near 80 nGy/hr and another near 200 nGy/hr), it could indicate two distinct geological sources of radiation.

3. Detecting Outliers and Anomalies

Histograms help identify extreme values that **deviate** from the normal distribution.

- Outliers may indicate **errors in data collection** or **unusual environmental conditions**.

Example: If most radiation measurements are **below 150 nGy/hr**, but a few exceed **500 nGy/hr**, those locations may require **further investigation**.

4. Comparing Different Data Sets

Histograms allow easy comparison of multiple datasets by overlaying different distributions.

Example: Comparing radiation levels from different regions (e.g., **Istria vs. a nuclear power plant zone**) using histograms can reveal how safe or hazardous each region is.

5. Choosing the Right Bin Size Matters

The level of detail in a histogram depends on **bin width**:

- **Too few bins** → Oversimplifies the data, hiding important details.
- **Too many bins** → Creates noise, making interpretation harder.

Example: In the **Istrian Peninsula histograms**, using 5 bins smoothed the distribution, while 15 bins revealed **more granular details**.

6. Aids in Probability and Statistical Analysis

- Used to **approximate probability distributions** (e.g., normal, Poisson, or exponential).
- Helps estimate measures like **mean, variance, and standard deviation**.

Example: In radiation safety, understanding the probability of exceeding **safe exposure levels** is crucial. Histograms help estimate these probabilities.

Conclusion: Why Use Histograms?

- **Visualizes Data Easily** → Helps see patterns, trends, and outliers.
- **Simplifies Large Datasets** → Useful when dealing with thousands of data points.
- **Assists in Decision Making** → Identifies safe vs. hazardous conditions.
- **Improves Statistical Analysis** → Supports hypothesis testing and probability estimates.

1.5.10 Disadvantages of Histograms

While histograms are useful for visualizing data distributions, they have some limitations:

1. Bin Size Sensitivity

- The choice of **bin width** significantly affects the shape of the histogram.
- **Too few bins** → Oversimplifies the distribution (hides patterns).
- **Too many bins** → Adds unnecessary noise, making it hard to interpret.

Example: A histogram with **5 bins** may look smooth, while one with **50 bins** may appear jagged.

2. Loss of Individual Data Points

- Histograms **group data into bins**, meaning **individual data points are not visible**.
- This makes it **impossible to see exact values** or detect small variations.

Alternative: Use a **scatter plot or KDE (Kernel Density Estimation)** to preserve individual points.

3. Difficult to Compare Multiple Distributions

- Comparing multiple histograms requires **overlaid them**, which can lead to **cluttered visuals**.

Alternative: Use a **box plot or violin plot** for better side-by-side comparison.

4. Misleading Interpretations

- The histogram's appearance can be **misleading depending on the bin width or starting point.**
- Different bin choices can make the **same dataset appear skewed or normally distributed.**

Example: A shift in bin edges can change the **perceived peaks and valleys** in the data.

5. Not Ideal for Small Datasets

- Histograms are **less effective with small datasets** because binning may result in **empty or uninformative bins.**

Alternative: A **dot plot or strip plot** is better for small sample sizes.

6. Cannot Show Exact Values or Outliers Easily

- Unlike scatter plots or box plots, histograms **do not explicitly show outliers.**
- Extreme values may be **hidden** within wide bins.

Solution: Using a **box plot** to visualize outliers more clearly.

When to Use and Avoid Histograms

Use Histogram When...	Avoid Histogram When...
We need to visualize frequency distributions.	We need to see individual data points.
The dataset is large enough to justify binning.	The dataset is too small (bins may be empty).
We want to compare shape, skewness, and spread.	We need exact values or outlier detection.